

Chapter 2, Optimization

Newton's method with a large p

Ying Wei & Xiaoqi Lu

February 26, 2020

1 Newton's method with a large p

- Recall the optimization $\hat{\theta} = \arg \max_{\theta \in R^p} f(\theta)$
- Newton's method suggests to update θ iteratively, such that the i th step is given by

$$\theta_i = \theta_{i-1} - [\nabla^2 f(\theta_{i-1})]^{-1} \nabla f(\theta_{i-1}),$$

where $\nabla f(\theta_{i-1})$ is the gradient, and $\nabla^2 f(\theta_{i-1})$ is the Hessian matrix.

- Question: is Newton's method scalable with increasing number of parameters p ?
 - The computational burden in calculating the inverse of the Hessian Matrix $[\nabla^2 f(\theta_{i-1})]^{-1}$ increases quickly with p .

1.1 Quasi-Newton Methods

Ascent direction: For a function f , a direction \mathbf{d} is an ascent direction for f at a given point $\boldsymbol{\theta}_0$ if there exists some $\epsilon > 0$ such that

$$f(\boldsymbol{\theta}_0 + \lambda \mathbf{d}) > f(\boldsymbol{\theta}_0)$$

for all $0 < \lambda < \epsilon$.

Directional derivative: The derivative of a function $f : \mathbb{R}^p \rightarrow \mathbb{R}$ at $\boldsymbol{\theta}$ in the direction of \mathbf{d} is defined by

$$\lim_{\lambda \rightarrow 0} \frac{f(\boldsymbol{\theta} + \lambda \mathbf{d}) - f(\boldsymbol{\theta})}{\lambda} = \left. \frac{\partial}{\partial \lambda} f(\boldsymbol{\theta} + \lambda \mathbf{d}) \right|_{\lambda=0} = \mathbf{d}' \nabla f(\boldsymbol{\theta}).$$

From this definition, we can see that \mathbf{d} is an ascent direction for f at $\boldsymbol{\theta}_0$ if and only if $\mathbf{d}' \nabla f(\boldsymbol{\theta}_0) > 0$.

Newton's method updates the parameters by

$$\boldsymbol{\theta}_{i+1} = \boldsymbol{\theta}_i - [\nabla^2 f(\boldsymbol{\theta}_i)]^{-1} \nabla f(\boldsymbol{\theta}_i),$$

Newton's direction $\mathbf{d} = -[\nabla^2 f(\boldsymbol{\theta}_i)]^{-1} \nabla f(\boldsymbol{\theta}_i)$ is an ascent direction if $[\nabla^2 f(\boldsymbol{\theta}_i)]^{-1}$ is negative definite.

More general: One can update

$$\boldsymbol{\theta}_{i+1} = \boldsymbol{\theta}_i + \mathbf{H}_{i,p \times p} \nabla f(\boldsymbol{\theta}_i),$$

and $f(\boldsymbol{\theta}_{i+1}) > f(\boldsymbol{\theta}_i)$ for any $\mathbf{H}_{i,p \times p}$ that is positive definite.

- Gradient Descent Algorithm: $\mathbf{H}_i = I_{p \times p}$ for any i
- Easy to compute, but could slow in convergence.

1.1.1 Other Quasi-Newton Iterations

$$\boldsymbol{\theta}_{i+1} = \boldsymbol{\theta}_i - \lambda_i [\mathbf{B}_i]^{-1} \nabla f(\boldsymbol{\theta}_i)$$

Question: can we find a surrogate matrix \mathbf{B}_i that is similar to $\nabla^2 f(\boldsymbol{\theta}_i)$, but easier to compute?

- Consider Taylor expansion on $\nabla f(\boldsymbol{\theta}_{i-1})$ around $\boldsymbol{\theta}_i$:

$$\nabla f(\boldsymbol{\theta}_i) - \nabla f(\boldsymbol{\theta}_{i-1}) = \nabla^2 f(\boldsymbol{\theta}_i) \cdot (\boldsymbol{\theta}_i - \boldsymbol{\theta}_{i-1}) + o(\|\boldsymbol{\theta}_i - \boldsymbol{\theta}_{i-1}\|) \quad (1)$$

- A good approximation can be achieved \mathbf{B}_i satisfies the following **secant equation**

$$\mathbf{B}_i \mathbf{S}_{i-1} = \mathbf{Y}_{i-1} \quad (2)$$

where

$$\mathbf{Y}_{i-1} = \nabla f(\boldsymbol{\theta}_i) - \nabla f(\boldsymbol{\theta}_{i-1})$$

$$\mathbf{S}_{i-1} = (\boldsymbol{\theta}_i - \boldsymbol{\theta}_{i-1}).$$

- There are infinitely many \mathbf{B}_i satisfying the **secant equation**.

SR1 (Symmetric-Rank-1) Method

Assuming a simple structural constrain that

$$\mathbf{B}_i = \mathbf{B}_{i-1} + \sigma v v^T,$$

where v is a p -dimensional vector. Combined with the secant equation, we have

$$\mathbf{Y}_{i-1} = \mathbf{B}_{i-1} \mathbf{S}_{i-1} + \sigma v^T \mathbf{S}_{i-1} v$$

$$\Rightarrow v = \delta (\mathbf{Y}_{i-1} - \mathbf{B}_{i-1} \mathbf{S}_{i-1}) \quad \text{for some } \delta \in \mathbb{R}$$

$$\Rightarrow (\mathbf{Y}_{i-1} - \mathbf{B}_{i-1} \mathbf{S}_{i-1}) = \sigma \delta^2 [\mathbf{S}_{i-1}^T (\mathbf{Y}_{i-1} - \mathbf{B}_{i-1} \mathbf{S}_{i-1})] (\mathbf{Y}_{i-1} - \mathbf{B}_{i-1} \mathbf{S}_{i-1})$$

Choose $\sigma = \text{sign} [\mathbf{S}_{i-1}^T (\mathbf{Y}_{i-1} - \mathbf{B}_{i-1} \mathbf{S}_{i-1})]$ and

$\delta = |\mathbf{S}_{i-1}^T (\mathbf{Y}_{i-1} - \mathbf{B}_{i-1} \mathbf{S}_{i-1})|^{-1/2}$, thus

$$\mathbf{B}_i = \mathbf{B}_{i-1} + \frac{(\mathbf{Y}_{i-1} - \mathbf{B}_{i-1} \mathbf{S}_{i-1})(\mathbf{Y}_{i-1} - \mathbf{B}_{i-1} \mathbf{S}_{i-1})^T}{\mathbf{S}_{i-1}^T (\mathbf{Y}_{i-1} - \mathbf{B}_{i-1} \mathbf{S}_{i-1})} \quad (3)$$

DFP (Davidon-Fletcher-Powell) Method

DFP Update

$$\min_{\mathbf{B}} \quad \|\mathbf{B} - \mathbf{B}_{i-1}\| \quad (4)$$

$$\text{subject to} \quad \mathbf{B} = \mathbf{B}^T, \mathbf{B}\mathbf{S}_{i-1} = \mathbf{Y}_{i-1} \quad (5)$$

Solution:

$$\mathbf{B}_i = \left(I - \frac{\mathbf{Y}_{i-1}\mathbf{S}_{i-1}^T}{\mathbf{Y}_{i-1}^T\mathbf{S}_{i-1}}\right)\mathbf{B}_k\left(I - \frac{\mathbf{S}_{i-1}\mathbf{Y}_{i-1}^T}{\mathbf{Y}_{i-1}^T\mathbf{S}_{i-1}}\right) + \frac{\mathbf{Y}_{i-1}\mathbf{Y}_{i-1}^T}{\mathbf{Y}_{i-1}^T\mathbf{S}_{i-1}} \quad (6)$$

BFGS (Broyden-Fletcher-Goldfarb-Shanno) Method*

To avoid taking inverse of \mathbf{B} , BFGS propose to approximate $\nabla^2 f(\boldsymbol{\theta}_i)^{-1}$ directly. Similar to DFP, they propose the following optimization problem:

$$\min_{\mathbf{H}} \quad \|\mathbf{H} - \mathbf{H}_{i-1}\| \quad (7)$$

$$\text{subject to} \quad \mathbf{H} = \mathbf{H}^T, \mathbf{H}\mathbf{Y}_{i-1} = \mathbf{S}_{i-1} \quad (8)$$

Solution:

$$\mathbf{H}_i = \left(I - \frac{\mathbf{S}_{i-1}\mathbf{Y}_{i-1}^T}{\mathbf{Y}_{i-1}^T\mathbf{S}_{i-1}}\right)\mathbf{H}_{i-1}\left(I - \frac{\mathbf{Y}_{i-1}\mathbf{S}_{i-1}^T}{\mathbf{Y}_{i-1}^T\mathbf{S}_{i-1}}\right) + \frac{\mathbf{S}_{i-1}\mathbf{S}_{i-1}^T}{\mathbf{Y}_{i-1}^T\mathbf{S}_{i-1}} \quad (9)$$

Note: BFGS is more effective than most quasi-Newton methods, and is the “go-to” method in many optimization problems.

1.2 Coordinate-wise optimization

Another simple approach is to consider coordinate descent approach, that starts with initial guess of $\boldsymbol{\theta}^{(0)} = (\theta_1^{(0)}, \theta_2^{(0)}, \dots, \theta_p^{(0)})$, and then update one component of $\boldsymbol{\theta}^{(0)}$ at a time iteratively

$$\theta_1^{(i+1)} = \arg \max_{\theta_1} f(\boldsymbol{\theta}_1, \theta_2^{(i)}, \dots, \theta_p^{(i)})$$

$$\theta_2^{(i+1)} = \arg \max_{\theta_2} f(\theta_1^{(i+1)}, \boldsymbol{\theta}_2, \theta_3^{(i)}, \dots, \theta_p^{(0)})$$

$$\theta_3^{(i+1)} = \arg \max_{\theta_3} f(\theta_1^{(i+1)}, \theta_2^{(i+1)}, \boldsymbol{\theta}_3, \theta_4^{(0)}, \dots, \theta_p^{(0)})$$

$$\vdots = \vdots$$

$$\theta_p^{(i+1)} = \arg \max_{\theta_p} f(\theta_1^{(i+1)}, \theta_2^{(i+1)}, \dots, \theta_{p-1}^{(i+1)}, \boldsymbol{\theta}_p)$$

Question: are we able to reach the global optimizer by maximizing / minimizing along each coordinate axis?

Answers: Yes, if $f(\boldsymbol{\theta})$ is convex and differentiable; $f(\boldsymbol{\theta})$ is convex but not differentiable; the coordinate descent approach still works if there exist a convex and differentiable $g(\cdot)$ and convex $h_k(\theta_k)$ for each $k = 1, \dots, p$ such that $f(\boldsymbol{\theta}) = g(\boldsymbol{\theta}) + \sum_{k=1}^p h_k(x_k)$.

- Solving p one-dimensional optimization is easier than solving one p -dimensional optimization.
- Order of cycle through coordinates is arbitrary
- One can replace individual coordinates with blocks of coordinates

Exercise 1.1 Consider a linear regression $EY = \mathbf{X}^T \boldsymbol{\beta}$, where Y is response vector, $\boldsymbol{\beta}$ is p -dimensional coefficient, and \mathbf{X} is the design matrix with columns $\mathbf{X}_1, \dots, \mathbf{X}_p$. The LS loss function

$$f(\boldsymbol{\beta}) = \|Y - \mathbf{X}^T \boldsymbol{\beta}\|^2$$

Consider minimizing over β_k while fixing all $\beta_j, j \neq k$

$$0 = \nabla_k f(\boldsymbol{\beta}) = \mathbf{X}_k^T (\mathbf{X}^T \boldsymbol{\beta} - Y) = \mathbf{X}_k^T (\mathbf{X}_k \beta_k + \mathbf{X}_{-k}^T \boldsymbol{\beta}_{-k} - Y)$$

$$\Rightarrow \beta_k = \frac{\mathbf{X}_k^T (Y - \mathbf{X}_{-k}^T \boldsymbol{\beta}_{-k})}{\mathbf{X}_k^T \mathbf{X}_k}$$

Coordinate descent repeats this update for $k = 1, 2, \dots, p, 1, 2, \dots$

1.3 Regularized regressions for high-dimensional p

Regularization is the common variable selection approaches for high-dimensional covariates. The best known Regularization is called LASSO. In linear regression, LASSO minimize

$$f(\beta) = \frac{1}{2} \sum_{i=1}^n (y_i - \sum_{j=1}^p x_{i,j} \beta_j)^2 + \gamma \sum_{j=1}^p |\beta_j|$$

for some $\gamma \geq 0$. Here the $x_{i,j}$ are standardized so that $\sum_i x_{i,j}/n = 0$ and $\sum_i x_{i,j}^2 = 1$.

With a single predictor x , the lasso solution is very simple

$$\hat{\beta}^{lasso}(\gamma) = S(\hat{\beta}, \gamma) = \text{sign}(\hat{\beta})(|\hat{\beta}| - \gamma)_+$$

$$\hat{\beta}^{\text{lasso}}(\gamma) = S(\hat{\beta}, \gamma) = \begin{cases} \hat{\beta} - \gamma, & \text{if } \hat{\beta} > 0 \text{ and } \gamma < |\hat{\beta}| \\ \hat{\beta} + \gamma, & \text{if } \hat{\beta} < 0 \text{ and } \gamma < |\hat{\beta}| \\ 0, & \text{if } \gamma > |\hat{\beta}| \end{cases}$$

- $S(\hat{\beta}, \gamma)$ is called soft threshold.
- If x are not standardized, i.e. $\langle x, x \rangle = \sum_i x_i^2 \neq 1$

$$\hat{\beta}^{\text{lasso}}(\gamma) = \frac{S(\langle x, y \rangle, \gamma)}{\langle x, x \rangle} = S(\hat{\beta}, \frac{\gamma}{\langle x, x \rangle})$$

- If multiple predictors that are uncorrelated/orthogonal (i.e. $\langle X_i, X_j \rangle = 0$), the lasso solutions are soft-thresholded versions of the individual least squares estimates.
- That is not the case when predictors are correlated. When p is large, the optimization could be challenging.

A coordinate-wise descent algorithm

- Coordinate-wise objective function

$$f(\beta_j) = \frac{1}{2} \sum_{i=1}^n (y_i - \sum_{k \neq j} x_{i,k} \tilde{\beta}_k - x_{i,j} \beta_j)^2 + \gamma \sum_{k \neq j} |\tilde{\beta}_k| + \gamma |\beta_j|$$

- Minimizing $f(\beta_j)$ w.r.t. β_j while having $\tilde{\beta}_k$ fixed, we have

$$\tilde{\beta}_j(\gamma) \leftarrow S \left(\sum_{i=1}^n x_{i,j} (y_i - \tilde{y}_i^{(-j)}), \gamma \right) \quad (10)$$

where $\tilde{y}_i^{(-j)} = \sum_{k \neq j} x_{i,k} \tilde{\beta}_k$. That is equivalent to regressing the partial residual $y_i - \tilde{y}_i^{(-j)}$ against $x_{i,j}$. The soft-threshold holds.

- We can then update β_j repeatedly for $j = 1, 2, \dots, p, 1, 2, \dots$ until convergence.

Covariance Updates and its flexibility with sparse matrix Note that $y_i - \tilde{y}_i^{(-j)} = y_i - \hat{y}_i + x_{i,j}\tilde{\beta}_j = r_i + x_{i,j}\tilde{\beta}_j$, where \hat{y}_i is fitted value at “current parameter” and r_i is the current residual.

$$\begin{aligned} \frac{1}{n} \sum_{i=1}^n x_{i,j} (y_i - \tilde{y}_i^{(-j)}) &= \frac{1}{n} \sum_{i=1}^n x_{i,j} r_i + \tilde{\beta}_j \\ &= \frac{1}{n} \left\{ \langle x_j, y \rangle - \sum_{k: |\tilde{\beta}_k| > 0} \langle x_j, x_k \rangle \tilde{\beta}_k \right\} + \tilde{\beta}_j \end{aligned}$$

where the inner product $\langle x_j, y \rangle = \sum_{i=1}^n x_{i,j} y_i$.

Sparse coding is an efficient way to store large sparse matrix, where we store only the non-zero entries and the coordinates where they occur.

Weighted Updates

- Often a weight w_i is associated with each observation.
- In this case, the regression coefficients is equivalent to regress $\sqrt{w_i}y_i$ against $\sqrt{w_i}x_i$

$$\sum_i w_i (y_i - x_i^\top \beta)^2 \Leftrightarrow \sum_i (\sqrt{w_i}y_i - \sqrt{w_i}x_i^\top \beta)^2$$

- The lasso update becomes only slightly more complicated:

$$\tilde{\beta}_j(\gamma) \leftarrow \frac{S\left(\sum_i w_i x_{i,j} (y_i - \tilde{y}_i^{(-j)}), \gamma\right)}{\sum_i w_i x_{i,j}^2} \quad (11)$$

Pathwise coordinatewise optimization algorithms

1. Starting at the smallest value λ for which the entire vector $\hat{\beta} = 0$.
 $\lambda_{\max} = \max_l \langle X_l, y \rangle$
2. Compute the solution sequentially at sequence
 $\lambda_{\max} \geq \lambda_1 \geq \cdots \geq \lambda_{\min} \geq 0$
3. For tuning parameter value λ_{k+1} , initialize coordinate descent algorithm at the computed solution for λ_k (warm start)

1.3.1 Extension of CD to logistic regression

Recall that the log likelihood of a logistic regression

$$f(\beta_0, \beta_1) = \sum_{i=1}^n \left(y_i(\beta_0 + \beta_1^T \mathbf{x}_i) - \log \left(1 + e^{\beta_0 + \beta_1^T \mathbf{x}_i} \right) \right).$$

The gradient of this function is

$$\nabla f(\beta_0, \beta_1) = \begin{pmatrix} \sum_{i=1}^n y_i - p_i \\ \sum_{i=1}^n \mathbf{x}_i (y_i - p_i) \end{pmatrix}_{(p+1) \times 1}, \quad (12)$$

where $p_i = P(Y_i = 1 | \mathbf{x}_i) = \frac{\exp(\beta_0 + \mathbf{x}_i^T \beta_1)}{1 + \exp(\beta_0 + \mathbf{x}_i^T \beta_1)}$.

The Hessian is given by

$$\begin{aligned} \nabla^2 f(\beta_0, \beta_1) &= - \sum_{i=1}^n \begin{pmatrix} 1 \\ \mathbf{x}_i \end{pmatrix} \begin{pmatrix} 1 & \mathbf{x}_i^T \end{pmatrix} p_i (1 - p_i) \\ &= - \begin{pmatrix} \sum p_i (1 - p_i) & \sum \mathbf{x}_i^T p_i (1 - p_i) \\ \sum \mathbf{x}_i p_i (1 - p_i) & \sum \mathbf{x}_i \mathbf{x}_i^T p_i (1 - p_i) \end{pmatrix}. \end{aligned}$$

A quadratic approximation to the log-likelihood $f(\beta_0, \beta_1)$

If we Taylor expansion the log-likelihood around “current estimates” $(\tilde{\beta}_0, \tilde{\beta}_1)$, we have

$$f(\beta_0, \beta_1) \approx \ell(\beta_0, \beta_1) = -\frac{1}{2n} \sum_{i=1}^n w_i (z_i - \beta_0 - \mathbf{x}_i^T \beta_1)^2 + C(\tilde{\beta}_0, \tilde{\beta}_1)$$

where

$$z_i = \tilde{\beta}_0 + \mathbf{x}_i^T \tilde{\beta}_1 + \frac{y_i - \tilde{p}_i(\mathbf{x}_i)}{\tilde{p}_i(\mathbf{x}_i)(1 - \tilde{p}_i(\mathbf{x}_i))} \text{ working response}$$

$$w_i = \tilde{p}_i(\mathbf{x}_i)(1 - \tilde{p}_i(\mathbf{x}_i)), \text{ working weights}$$

$$\tilde{p}_i = \frac{\exp(\tilde{\beta}_0 + \mathbf{x}_i^T \tilde{\beta}_1)}{1 + \exp(\tilde{\beta}_0 + \mathbf{x}_i^T \tilde{\beta}_1)} \text{ (probability evaluated at the current parameters)}$$

Exercise 1.2 *The logistic-lasso can be written as a penalized weighted least-squares problem*

$$\min_{(\beta_0, \boldsymbol{\beta}_1)} L(\beta_0, \boldsymbol{\beta}_1, \lambda) = \{-\ell(\beta_0, \boldsymbol{\beta}_1) + \lambda \sum_{j=0}^p |\beta_j|\}$$

Derive path-wise coordinate descent algorithm update for the optimization above

Step 1 *Find λ_{max} such that all the estimated β are zero;*

Step 2 *Define a fine sequence $\lambda_{max} \geq \lambda_1 \geq \cdots \geq \lambda_{min} \geq 0$*

Step 3 *Defined the quadratic approximated objective function*

*$L(\beta_0, \boldsymbol{\beta}_1, \lambda)$ for λ_k using the estimated parameter at λ_{k-1}
($\lambda_{k-1} > \lambda_k$).*

Step 4 *Run coordinate descent algorithm to find the optimization defined in Step 3*